

Transfer Learning fundamentals

1) Concept

The idea of **transfer learning** is using a **pre-trained** model (a model that has been trained on a large dataset) from outside sources, **implement** it a bit, and apply on your problems. As a reminder, the early layers learn more abstract features than the later ones. Therefore, when using transfer learning, we keep the early ones and modify the later one. Due to this fact can apply transfer learning on classes that were not trained on the pre-trained model. I will use a pre-trained convolutional network as an illustration for my transfer learning practice today.

A convolutional network consists of conv layers and dense layers. The conv layers in the pre-trained model are called **conv base**. Implementing a pre-trained model means replacing your dense layers and train the model in a way that does not affect the conv base weights too much. There are two popular implement methods: **feature extraction** and **fine tune**.

- Feature extraction: Only modify dense layers
- Fine-tune: modify the last few conv base layers

In the code below, I will use the fine-tune approach that involves tweaking the last conv layer

The goal is to use the pre-trained model **VGG19** to classify **Santa** or **not Santa**

2 Data and VGG19

The data includes 385 santa images and 447 not santa images that are divided into 3 folder

- Train: 271 Images each
- Validation: 100 Images each
- Test: 90 Images each

```
In [87]: train_folder = 'split/train'
test_folder = 'split/test'
val_folder = 'split/validation'
```

```
In [103]: from IPython.display import Image
from IPython.display import display
x = Image(filename=train_folder+'/santa/00000011.jpg',width=300)
y = Image(filename=train_folder+'/not_santa/00000001.jpg',width=300)
print("Santa")
display(x)
```

Santa



```
In [104]: print("Not Santa")
display(y)
```

Not Santa



VGG-19 is a convolutional neural network that is 19 layers deep. You can load a pretrained version of the network trained on more than a million images from the ImageNet database [1]. The pretrained network can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. (has no santa)

```
In [86]: from keras.applications import VGG19

#Exclude the dense layers and reshape from (224, 224 ,3) to (64, 64, 3)
cnn_base = VGG19(weights='imagenet',
                  include_top=False,
                  input_shape=(64, 64, 3))
cnn_base.summary()
```

Model: "vgg19"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	(None, 64, 64, 3)	0
block1_conv1 (Conv2D)	(None, 64, 64, 64)	1792
block1_conv2 (Conv2D)	(None, 64, 64, 64)	36928
block1_pool (MaxPooling2D)	(None, 32, 32, 64)	0
block2_conv1 (Conv2D)	(None, 32, 32, 128)	73856
block2_conv2 (Conv2D)	(None, 32, 32, 128)	147584
block2_pool (MaxPooling2D)	(None, 16, 16, 128)	0
block3_conv1 (Conv2D)	(None, 16, 16, 256)	295168
block3_conv2 (Conv2D)	(None, 16, 16, 256)	590080
block3_conv3 (Conv2D)	(None, 16, 16, 256)	590080
block3_conv4 (Conv2D)	(None, 16, 16, 256)	590080
block3_pool (MaxPooling2D)	(None, 8, 8, 256)	0
block4_conv1 (Conv2D)	(None, 8, 8, 512)	1180160
block4_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block4_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block4_conv4 (Conv2D)	(None, 8, 8, 512)	2359808
block4_pool (MaxPooling2D)	(None, 4, 4, 512)	0
block5_conv1 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv4 (Conv2D)	(None, 4, 4, 512)	2359808
block5_pool (MaxPooling2D)	(None, 2, 2, 512)	0
=====		
Total params: 20,024,384		
Trainable params: 20,024,384		
Non-trainable params: 0		

3) Action

```
In [78]: # Add the dense layers to conv base
from keras import models
from keras import layers
from keras import optimizers

model = models.Sequential() #Sequential model
model.add(cnn_base)
model.add(layers.Flatten())
model.add(layers.Dense(132, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=2e-5),
              metrics=['acc'])
```

```
In [79]: model.summary()

Model: "sequential_11"

Layer (type)                Output Shape                Param #
=====
vgg19 (Model)                (None, 2, 2, 512)          20024384
-----
flatten_6 (Flatten)         (None, 2048)                0
-----
dense_18 (Dense)             (None, 132)                 270468
-----
dense_19 (Dense)             (None, 1)                   133
=====
Total params: 20,294,985
Trainable params: 270,601
Non-trainable params: 20,024,384
=====
```

```
In [81]: #Perform Data Augmentation with Image Data Generator
# Step 1: Define ImageDataGenerator object with different parameters
# Step 2: REPLACE the original image with 20 different images through data aug by flow_from_directory
# Note that we set "test set batch size" = "test set size"

from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale=1./255,
                                   rotation_range=40,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   horizontal_flip=True,
                                   fill_mode='nearest')

train_generator = train_datagen.flow_from_directory(train_folder,
                                                    target_size=(64, 64),
                                                    batch_size= 20,
                                                    class_mode= 'binary')

val_generator = ImageDataGenerator(rescale=1./255).flow_from_directory(val_folder,
                                                                        target_size=(64, 64),
                                                                        batch_size=20,
                                                                        class_mode='binary')

test_generator = ImageDataGenerator(rescale=1./255).flow_from_directory(test_folder,
                                                                        target_size=(64, 64),
                                                                        batch_size=180,
                                                                        class_mode='binary')

test_images, test_labels = next(test_generator)
```

Found 542 images belonging to 2 classes.
Found 200 images belonging to 2 classes.
Found 180 images belonging to 2 classes.

```
In [82]: #Fine tune

#Step 1: Free cnn base
#Step 2: unfreeze the last conv layers in cnn

cnn_base.trainable = False
cnn_base.block5=True
for layer in model.layers:
    print(layer.name, layer.trainable)
```

vgg19 False
flatten_6 True
dense_18 True
dense_19 True

```
In [83]: history = model.fit_generator(train_generator,
                                     epochs=10,
                                     validation_data=val_generator,
                                     validation_steps=10)
```

Epoch 1/10
28/28 [=====] - 29s 1s/step - loss: 0.6654 - acc: 0.6015 - val_loss: 0.6127 - val_acc: 0.6850
Epoch 2/10
28/28 [=====] - 32s 1s/step - loss: 0.6110 - acc: 0.7103 - val_loss: 0.6224 - val_acc: 0.7700
Epoch 3/10
28/28 [=====] - 33s 1s/step - loss: 0.5678 - acc: 0.7823 - val_loss: 0.4688 - val_acc: 0.8350
Epoch 4/10
28/28 [=====] - 33s 1s/step - loss: 0.5257 - acc: 0.8173 - val_loss: 0.5246 - val_acc: 0.8300
Epoch 5/10
28/28 [=====] - 33s 1s/step - loss: 0.4962 - acc: 0.8432 - val_loss: 0.4710 - val_acc: 0.9100
Epoch 6/10
28/28 [=====] - 33s 1s/step - loss: 0.4630 - acc: 0.8542 - val_loss: 0.3898 - val_acc: 0.8950
Epoch 7/10
28/28 [=====] - 34s 1s/step - loss: 0.4470 - acc: 0.8653 - val_loss: 0.3641 - val_acc: 0.9150
Epoch 8/10
28/28 [=====] - 36s 1s/step - loss: 0.4277 - acc: 0.8653 - val_loss: 0.3535 - val_acc: 0.9150
Epoch 9/10
28/28 [=====] - 55s 2s/step - loss: 0.4087 - acc: 0.8672 - val_loss: 0.3937 - val_acc: 0.9150
Epoch 10/10
28/28 [=====] - 37s 1s/step - loss: 0.3930 - acc: 0.8672 - val_loss: 0.2952 - val_acc: 0.9150

```
In [91]: # Accuracy on the test set
test_loss, test_acc = model.evaluate_generator(test_generator, steps=50)
print('test acc:', test_acc)
```

test acc: 0.8944444465637207

Some discovery

Lemma: Sigmoid is a special case of softmax function

Softmax for K classes:

$$\Pr(Y_i = k) = \frac{e^{\beta_k \cdot X_i}}{\sum_{0 \leq c \leq K} e^{\beta_c \cdot X_i}}$$

where β_k is a vector of weights (or regression coefficients) corresponding to outcome k. When K=2, we have

$$\begin{aligned} \Pr(Y_i = 0) &= \frac{e^{\beta_0 \cdot X_i}}{\sum_{0 \leq c \leq K} e^{\beta_c \cdot X_i}} = \frac{e^{\beta_0 \cdot X_i}}{e^{\beta_0 \cdot X_i} + e^{\beta_1 \cdot X_i}} = \frac{e^{(\beta_0 - \beta_1) \cdot X_i}}{e^{(\beta_0 - \beta_1) \cdot X_i} + 1} = \frac{e^{-\beta \cdot X_i}}{1 + e^{-\beta \cdot X_i}} \\ \Pr(Y_i = 1) &= \frac{e^{\beta_1 \cdot X_i}}{\sum_{0 \leq c \leq K} e^{\beta_c \cdot X_i}} = \frac{e^{\beta_1 \cdot X_i}}{e^{\beta_0 \cdot X_i} + e^{\beta_1 \cdot X_i}} = \frac{1}{e^{(\beta_0 - \beta_1) \cdot X_i} + 1} = \frac{1}{1 + e^{-\beta \cdot X_i}} \end{aligned} \quad \text{Q.E.D}$$

Therefore, for a binary classification problem, in the output layer of a neural network, you can have either have a softmax with 2 units, or a sigmoid with 1 unit.